Analysis of Algorithms - Home Work 3

CSE 548 Fall '19

Prof. Jie Gao

Submission By:

Narayan Acharya

112734365

Contents

1   Question 1

Textbook [Kleinberg & Tardos] Chapter 4, page 190, problem #6.
Solution:

Let us say that the i^{th} contestant has a swimming, biking and running time denoted by s_i, b_i and r_i respectively. We are to order the contestants such that the last contestant to finish has the earliest possible finishing time. We also have a constraint that no 2 contestants can have overlapping swimming times. Although, there is NO such constraint on their biking and running times. Let us consider the biking and running time of i^{th} contestant to be p_i where $p_i = b_i + r_i$.

Consider the following way of ordering the contestants to achieve our goal: We order them in decreasing order of p_i. That is, for contestants i and j, if i has a higher p_i then he/she goes before p_j. We can argue that this is the most optimal way of ordering contestants by considering existence of another solution that does not follow above rule and is optimal.

Consider that there exists an optimal solution different from the one proposed above such that for contestants i and j, where i starts before j and $p_i < p_j$. Given that i start before j, i will leave the pool earlier and. Time to complete the race from when i started:

$$t_{i1} = s_i + p_i$$

$$t_{j1} = s_i + s_j + p_j$$

Comparing the two, we can easily see that i will finish earlier (strike our s_i from both, $p_i < p_j$ and $s_j > 0$). Thus, $t_{i1} < t_{j1}$

Let us swap i and j now and compare the times to complete from when j started:

$$t_{j2} = s_j + p_j$$

$$t_{i2} = s_j + s_i + p_i$$

Comparing these with t_{i1} and t_{j1}, we can easily see that:

$$t_{i1} < t_{j2} < t_{j1}$$

$$t_{i1} < t_{i2} < t_{j1}$$

In all cases t_{j1} is highest. That means our approach with the first option is not as good compared to our second options where we switched contestants and brought them in accordance of our proposed solution. Hence, it is better if we swap i and j.

2   Question 2

Textbook [Kleinberg & Tardos] Chapter 4, page 190, problem #12.
Solution:

(a) Our constraint is such that any time t, the total number of bits we've sent through from time
0 to t cannot exceed rt. The constraint does NOT bound number of bits we can send at any
specific time t. So, at any time we can send more bits than the rate specified, if we've sent
lesser bits in the previous stream.

So the constraint, $b_i < rt_i$ only holds for the first stream of bits. Subsequent streams may or
may not hold to this inequality. In either case they will still be valid. So the claim is false.

(b) We can use the 'Greedy Algorithm Stays Ahead' argument to prove our case. In order to find
a valid and feasible schedule we sort by bits we need to transmit in each stream i.e. We sort
by b_i/t_i. In such an ordering we always transmit the least number of bits we can possibly
send at that time. If at any time, our constraint does not hold true then there is no way
we can send another stream of bits ahead of this as it will at least transmit as many bits as
the current stream because we have sorted them in increasing order of bit-rate. The running
time of this algorithm will be bound by $O(nlogn)$, the amount of time we need to sort the
streams in increasing order of bit-rate.

To only find whether there exists a schedule where we can order all streams without breaking
the constraint we simply keep of track of number of bits 'saved' at each transmission. Bits
saved can be used to transmit another stream having a higher bit-rate than the allowed
maximum rate of r. In these cases, we can consider the bits saved to be negative. If the total
sum of all such bits saved is greater than 0, then we can claim that we can order streams in
some order without breaking the constraint. The running time for this will be $O(n)$ as we
only need to calculate bits saved at each bit-stream and keep track of total bits saved.

3   Question 3

Textbook [Kleinberg & Tardos] Chapter 4, page 190, problem #14.
Solution:

(a) We wish to minimize the number of times we invoke the script to check our critical processes. The greedy approach to this would be holding off on the check as long as possible and doing one if absolutely necessary so that we don't miss a process. One approach would be to sort the processes according to their finish times and to invoke the status check right before the process finishes if the process has not been checked at least once prior. Let us consider the correctness and optimality of our approach.

(1) Correctness: The approach ensures that all the processes are checked because we invoke the status check right before the end time of every process if it has not yet been checked. Thus it computes a valid set of status check invocation timings.

(2) Optimality: To argue that this approach computes the minimum number of status checks we use 'Greedy Algorithm Stays Ahead' argument. Consider, that our algorithm produces k status checks up to a certain process evaluation, and the set is denoted by S. Also, there exists another set of status checks, S'. We would like to show that S' too has at least k status checks up until the same process is evaluated. We prove this by induction:

   i. For $n = 1$, both S and S' will have to have at least 1 status check.
   ii. For some $n = m$, say S and S' have k status checks.
   iii. For $n = m+1$, we have two options, $(m+1)^{th}$ overlaps or does not overlap with our current process. In case it does, our k^{th} status check in S will take care of the status check for $(m+1)^{th}$ process. In case it does not, we will need to add our $(k+1)^{th}$ status check at the finish time of the $(m+1)^{th}$ process. In both cases, S' too will need to add at least one more check to cover the $(m+1)^{th}$ process. Thus we have $(k+1)$ checks for in both S and S'.

Thus our algorithm will perform at least as well any other optimal algorithm. The algorithm is bound by $O(nlogn)$ - the amount it would required to sort the processes by their finishing times and $O(n)$ - the amount of time to iterate over the sorted processes to insert the process checks. Thus, running time is bound by $O(nlogn)$.

(b)

4 Question 4

Textbook [Kleinberg & Tardos] Chapter 4, page 190, problem #16.
Solution:

5   Question 5

Textbook [Kleinberg & Tardos] Chapter 5, page 246, problem #1.
Solution:

6   Question 6

Textbook [Kleinberg & Tardos] Chapter 5, page 246, problem #7.
Solution:

References

[1] ,